

Notes (in progress) about P2P Web

Rasmus Erik Voel Jensen

2017

Abstract

Random ramblings and notes during the P2P Web.

General notes

Purpose:

- Infrastructure for no-server HTML5 apps => a decentralized trustless computer for the web

In short:

- Network topology: kademlia like, - address = hash of pubkey
- State/storage: each node stores a neighbourhood around its own address, saved in blockchain merkel tree
- Operations: changes to state are verifiable, and verified by nodes in neighbourhood
- Balance: nodes get paid for doing tasks for the network, and can use this to buy tasks in the network. Also pay/payout for state blockchain.
- Tasks: stored, nodes assigned to tasks in deterministic random part of storage, proof-of-result stored, result stored, verification/value, balance updated.

Additional notes:

- WebPlatform: computations in webassembly. WebRTC as transport (thus modified kademlia). Crypto-algorithms from crypto.subtle.
- Neighbourhood size and amount of state per node - determined by node density (global minimum density / local density). Fixed amount of memory per node.
- Mutable references in blockchain (using balance to keep alive)
- Autonomous processes (using balance to keep alive)
- Not entire blockchain stored, only parts needed by the node

- Stake in computation tasks
- Balance/trade between processes
- Introduced ‘errors’ in blockchain, and bounties for finding/proving them.
- Binary/Quad merkel tree for proofs
- Pub/private-key derived from entropy source
- Task types: computation, storage, storage-transfer, find node with certain data
- Node trust / reliability proof via blockchain
- Block-tree rather than chain
- Computational task level of validation
- Result safety: added to state by any node in neighbourhood by proof of distance of computing-node to task.
- Computational task: computing time bound, and cost calculation.
- consensus algorithm: CRDT, additional data in timeinterval: after last block, before timed signature from other deterministic random node
- Tagged overlay network - opt-in part of infrastructure for tunable bandwidth requirements
- Network simulation (core optimised for low memory)
- Bandwidth optimised, - number of significant bits per node-id, stream compression, only send diffs etc.

Explore/ideas:

- Performance characteristics of current WebRTC implementations
- Performance effects of design choices for Kademlia-like algorithm on top of WebRTC (instead of UDP)
- Verifiable “computational” tasks, and economy based “computation”.
- (Survey p2p overlay networks)
- WebRTC bootstrapping options (decentralised signalling server vs. actual node)
- Infrastructure deployment - bootstrap-code + load signed version of code from network, - partly test within network before full deploy.

Description of algorithm:

- nodes connected in kademlia-like structure
- regular state snapshot (blockchain merkel-dag)
 - divide-and-conquer consensus algorithm, verifying credit updates in neighbourhood.
 - each node stores the state of a neighbourhood around its own address, as well as the path to the root. The neighbourhood size is fixed for all, ensuring good redundancy of data for
- content of state
 - list of entities(nodes)

- * id
- * balance/credits (updated by work, tasks, cost of staying in blockchain, and transfers)
- * state (+ proof)
- * tasks - scheduled for execution - wager
- * result of previous scheduled task
- * work
 - stake
 - result
 - proof-of-work
- * state
- verifiable tasks
- task types
 - computation
 - storage
 - * data
 - * key/value
 - random verifications (of proof-of-stake tasks)
 - blockchain verification
- entities
 - nodes
 - nodes with stake
 - accounts (pub-key)
 - autonomous
- computational process
 - task gets stored in blockchain
 - task gets assigned to a number of bcrandom nodes
 - task gets done, and proof-of-work gets stored in the blockchain
 - task result gets released
 - result+proof-of-work get validated + signed into blockchain
 - balance is updated

Design criteria

- low bandwidth
- low memory footprint (useful for large simulation, as well as embedded systems)
- low code footprint
- tagging of hosts
- connect to arbitrary host
- foundation for other p2p applications

Roadmap

TODO: describe here

notes

Ideas of places for publications:

- Open Access Journals (NB: <https://doaj.org>)
 - Ledger Journal
 - EAI Transactions on Scalable Information Systems
 - Computer Science (AGH)
- ACM/IEEE (pay for open access)
 - IEEE Transactions on Parallel and Distributed Systems
 - ACM Transactions on Computer Systems (TOCS)
- Conferences (in Europe)
 - 2017-10-30 2018-02-26/28 IFIP NTMS 1st International Workshop on Blockchains and Smart Contracts (BSC)
 - ~~2018-02-05 2018-06-20/22 DCAI 2018 : 15th International Conference on Distributed Computing and Artificial Intelligence~~
- Central non-conference sources (based on related articles)
 - “The Computing Research Repository” (not a Journal, but arXiv, indexed in dblp etc.)
 - Peer-to-Peer Networking and Applications (Springer)
 - ~~Computer Networks (avoid Elsevier)~~

Literature

Various articles etc. that I should take a look add (and possibly add notes here eventually)

- (Parno, Raykova, and Vaikuntanathan 2012)
- (Kaune et al. 2008)
- (Wood 2014)
- (Golem 2016)
- (Rhea et al. 2004)
- (Wang, Yang, and Chen 2010)
- (Benet 2014)
- (Maymounkov and Mazières 2002)
- (Gennaro, Gentry, and Parno 2010)
- (Medrano-Chávez, Pérez-Cortés, and Lopez-Guerrero 2015)
- (Ou et al. 2010)
- (Parno et al. 2013)

- (Baumgart and Mies 2007)
- (Jiménez, Osmani, and Knutsson 2011)
- (Surati, Jinwala, and Garg 2017)
- (Wood and Steiner 2016)
- (Haas et al. 2017)

<https://allquantor.at/blockchainbib/bibtex.html>

Notes for later / optimised version

- page size
- typically 4K (getpagesize() is 4K on my linux, and that looks like common size via https://en.wikipedia.org/wiki/Page_size)
- webassembly 64K page size
- minimise memory usage (for ability to run large simulations).
- i.e. 64K per nodes => 100K nodes in memory simulation ~ 6G memory

Bibliography

Baumgart, Ingmar, and Sebastian Mies. 2007. “S/Kademlia: A Practicable Approach Towards Secure Key-Based Routing.” In *13th International Conference on Parallel and Distributed Systems, ICPADS 2007, Hsinchu, Taiwan, December 5-7, 2007*, 1–8. doi:10.1109/ICPADS.2007.4447808.

Benet, Juan. 2014. “IPFS - Content Addressed, Versioned, P2P File System.” *CoRR* abs/1407.3561. <http://arxiv.org/abs/1407.3561>.

Gennaro, Rosario, Craig Gentry, and Bryan Parno. 2010. “Non-Interactive Verifiable Computing: Outsourcing Computation to Untrusted Workers.” In *Advances in Cryptology – Crypto 2010: 30th Annual Cryptology Conference, Santa Barbara, ca, Usa, August 15-19, 2010. Proceedings*, edited by Tal Rabin, 465–82. Berlin, Heidelberg: Springer Berlin Heidelberg. doi:10.1007/978-3-642-14623-7_25.

Golem. 2016. “The Golem Project - Crowdfunding Whitepaper.” <http://www.golemproject.net/doc/DraftGolemProjectWhitepaper.pdf>.

Haas, Andreas, Andreas Rossberg, Derek L. Schuff, Ben L. Titzer, Michael Holman, Dan Gohman, Luke Wagner, Alon Zakai, and J. F. Bastien. 2017. “Bringing the Web up to Speed with Webassembly.” In *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and*

Implementation, PLDI 2017, Barcelona, Spain, June 18-23, 2017, 185–200. doi:10.1145/3062341.3062363.

Jiménez, Raúl, Flutra Osmani, and Björn Knutsson. 2011. “Sub-Second Lookups on a Large-Scale Kademlia-Based Overlay.” In *2011 IEEE International Conference on Peer-to-Peer Computing, P2P 2011, Kyoto, Japan, August 31 - September 2, 2011*, 82–91. doi:10.1109/P2P.2011.6038665.

Kaune, S., T. Lauinger, A. Kovacevic, and K. Pussep. 2008. “Embracing the Peer Next Door: Proximity in Kademlia.” In *2008 Eighth International Conference on Peer-to-Peer Computing*, 343–50. doi:10.1109/P2P.2008.36.

Maymounkov, Petar, and David Mazières. 2002. “Kademlia: A Peer-to-Peer Information System Based on the XOR Metric.” doi:10.1007/3-540-45748-8_5.

Medrano-Chávez, Adán G., Elizabeth Pérez-Cortés, and Miguel Lopez-Guerrero. 2015. “A Performance Comparison of Chord and Kademlia Dhts in High Churn Scenarios.” *Peer-to-Peer Networking and Applications* 8 (5): 807–21. doi:10.1007/s12083-014-0294-y.

Ou, Zhonghong, Erkki Harjula, Otso Kassinen, and Mika Ylianttila. 2010. “Performance Evaluation of a Kademlia-Based Communication-Oriented P2P System Under Churn.” *Computer Networks* 54 (5): 689–705. doi:10.1016/j.comnet.2009.09.022.

Parno, Bryan, Jon Howell, Craig Gentry, and Mariana Raykova. 2013. “Pinocchio: Nearly Practical Verifiable Computation.” In *2013 IEEE Symposium on Security and Privacy, SP 2013, Berkeley, ca, Usa, May 19-22, 2013*, 238–52. doi:10.1109/SP.2013.47.

Parno, Bryan, Mariana Raykova, and Vinod Vaikuntanathan. 2012. “How to Delegate and Verify in Public: Verifiable Computation from Attribute-Based Encryption.” In *Theory of Cryptography: 9th Theory of Cryptography Conference, Tcc 2012, Taormina, Sicily, Italy, March 19-21, 2012. Proceedings*, edited by Ronald Cramer, 422–39. Berlin, Heidelberg: Springer Berlin Heidelberg. doi:10.1007/978-3-642-28914-9_24.

Rhea, Sean C., Dennis Geels, Timothy Roscoe, and John Kubiawicz. 2004. “Handling Churn in a DHT.” In *Proceedings of the General Track: 2004 USENIX Annual Technical Conference, June 27 - July 2, 2004, Boston Marriott Copley Place, Boston, Ma, USA*, 127–40. <http://www.usenix.org/publications/library/proceedings/usenix04/tech/general/rhea.html>.

Surati, Shivangi, Devesh C. Jinwala, and Sanjay Garg. 2017. “A Survey of Simulators for P2p Overlay Networks with a Case Study of the P2p Tree Overlay Using an Event-Driven Simulator.” *Engineering Science and Technology, an International Journal* 20 (2): 705–20.

doi:<http://dx.doi.org/10.1016/j.jestch.2016.12.010>.

Wang, C., N. Yang, and H. Chen. 2010. “Improving Lookup Performance Based on Kademlia.” In *2010 Second International Conference on Networks Security, Wireless Communications and Trusted Computing*, 1:446–49. doi:10.1109/NSWCTC.2010.111.

Wood, Gavin. 2014. “Ethereum: A Secure Decentralised Generalised Transaction Ledger.” *Ethereum Project Yellow Paper*. <https://github.com/ethereum/yellowpaper>.

Wood, Gavin, and Jutta Steiner. 2016. “Trustless Computing—The What Not the How.” In *Banking Beyond Banks and Money: A Guide to Banking Services in the Twenty-First Century*, edited by Paolo Tasca, Tomaso Aste, Lorian Pelizzon, and Nicolas Perony, 133–44. Cham: Springer International Publishing. doi:10.1007/978-3-319-42448-4_8.